Utkarsh Tyagi

📞 +1 240 501 3163   ✉ utkarsht@umd.edu   🌐 utkarsh4430.github.io   💼 Linkedin   🎓 Google Scholar

EDUCATION

University of Maryland (Advised by Prof. Dinesh Manocha)　　　　　**College Park, Maryland**
(Incoming CS PhD) M.S. in Computer Science, Advised by - GPA 3.85/4.0　　　*08/2023 – 05/2025*

Delhi Technological University (Formerly Delhi College of Engineering)　　　　　**Delhi, India**
*B.Tech in Computer Science - GPA **8.7**/10, Major GPA **9.05**/10*　　　*2017 – 2021*

PUBLICATIONS

MULTIMODAL LEARNING

- VDGD: Mitigating LVLM Hallucinations in Cognitive Prompts by Bridging the Visual Perception Gap
  NeurIPS 2024 (Under review)

- ProSE: Diffusion Priors for Speech Enhancement
  NeurIPS 2024 (Under review)

- CompA: Addressing the Gap in Compositional Reasoning in Audio-Language Models
  ICLR 2024

- Do Vision-Language Models Understand Compound Nouns?
  NAACL 2024

- LipGER: Visually-Conditioned Generative Error Correction for Robust Automatic Speech Recognition
  InterSpeech 2024

- AdVerb: Visually Guided Audio Dereverberation
  ICCV 2023

- MMER: Multimodal Multi-task Learning for Speech Emotion Recognition
  InterSpeech 2023

NATURAL LANGUAGE PROCESSING (NLP)

- ABEX: Data Augmentation for Low-Resource NLU via Expanding Abstract Descriptions
  ACL 2024

- ASPIRE: Language-Guided Data Augmentation for Improving Robustness Against Spurious Correlations
  ACL 2024

- CoDa: Constrained Generation based Data Augmentation for Low-Resource NLP
  NAACL 2024

- ACLM: Selective-Denoising based Data Augmentation for Low-Resource Complex NER
  ACL 2023

- DALE: Generative Data Augmentation for Low-Resource Legal NLP
  EMNLP 2023

- BioAug: Conditional Generation based Data Augmentation for Low-Resource Biomedical NER
  SIGIR 2023

INDUSTRY EXPERIENCE

Samsung SARC/ACL　　　　　**San Jose, California**
ML Research Intern - Efficient GenAI Inference　　　*May 2024 – Current*

- Optimizing state-of-the-art Large Language Models (LLMs) for deployment on resource-constrained devices.
- Finetuning LLMs and optimizing their performance through sophisticated Retrieval-Augmented Generation (RAG) techniques leveraging internal documentation databases.

Atlassian India LLP　　　　　**Bangalore, Karnataka**
Software Development Engineer 2　　　*July 2021 – August 2023*

- Worked on the JIRA Issue Create Experience to deliver a reliable 1yand performant interface, which helped improve key business metrics and customer satisfaction ratings.
- Achieved **99.996% frontend reliability** and scaled the experience for over 250k customers with more than 2M issues created every day. Improved performance TTI from 3.9s to **1.2s** using modern technology stack.

Software Engineer Intern *Jan 2021 – June 2021*
- Identified limitations and bottlenecks in the existing legacy API implementation and proposed comprehensive improvements for system revamping. Utilized this performant API to develop the user interface.
- Reduced the **latency** of modal load by **95.57%** by avoiding redundant server-side database queries and scaled to 1M customers.

Samsung R&D Institute Bangalore **Bangalore, Karnataka**
Machine Learning Intern *May 2020 – July 2020*
- Worked with the **Voice Intelligence** team of Samsung to improve Bixby's wakeup word detection.
- Researched keyword spotting techniques and built an efficient on-device LSTM-CTC based, vocabulary-independent keyword spotter

Research Experience

GAMMA Lab, University of Maryland **College Park, Maryland**
Research Assistant *August 2022 – Present*
- Currently researching multimodal learning; working on video representation learning for long-form video understanding
- Also focus on low-resource (labeled data and compute) learning with speech, NLP, or vision applications. In this area, I solve problems using self-supervised learning, synthetic data augmentation, etc.
- Published at **ICLR**, **ICCV**, **ACL**, **EMNLP**, **InterSpeech**, **SIGIR**, **NAACL**
- Advised by Prof. Dinesh Manocha

Multimodal Digital Media Analysis Lab **IIITD, Delhi**
Machine Learning Researcher *April 2022 – March 2023*
- Worked under the supervision of Dr. Rajiv Ratn Shah in the areas of Speech and Language Processing.
- Worked on multi-lingual automatic speech scoring systems for low-resource Indian languages.
- Also explored novel architectures to detect implicit hate speech in online conversations. Paper accepted at **AAAI 2023 Defactify**

Patents

- Apparatuses, methods, and computer program products for generating an abstract context summary scheduling interface configured for scheduling and outputting abstract context summaries for multi-party communication channels **U.S. Patent Application No. 17/936,695** (Patent Pending)

- Apparatuses, methods, and computer program products for generating and selectively outputting abstract context summaries for multi-party communication channels **U.S. Patent Application No. 17/936,705** (Patent Pending)

Technical Skills

Languages: (Highly Proficient) Python, JavaScript, C++ (Moderate) Java, SQL
Frameworks: Pytorch, FastAI, Tensorflow, Scikit-Learn, ReactJS, Jest, Enzyme, NLTK

Achievements

- **Reviewer** for ACL ARR 2024 (Feb, April, June), ICASSP 2024, EMNLP 2024 (Industry Track Reviewer)
- **People's choice award 2021 & 2022:** Atlassian's internal hackathon
- **Winner** of Digital Management Inc. India Hackathon, 2019
- **2nd position** in Student Hackday 2019 organised by Skillenza
- **Barclays India Hackathon 2019 Top 4**, All India
- Achieved **5 star rating** on **CodeChef** platform for competitive programming
- **JEE Mains 2017: 99.49%** percentile, Pan India